

Co-Processor Accelerated Data Management

Sebastian Breß



German
Research Center
for Artificial
Intelligence



sebastian.bress@dfki.de

Acknowledgements

These slides are based on the following material:

- Lecture Data Processing on GPUs and GPGPUs (René Müller, IBM Research)
- Lecture Data Processing on Modern Hardware (Jens Teubner, TU Dortmund)
- Tutorial Many-Core-Architekturen zur Datenbankbeschleunigung [14] (Kai-Uwe Sattler and Felix Beier, TU Ilmenau; Jens Teubner, TU Dortmund)

Thanks!

Performance Limitations of Modern Processors

Until 2004/2005:

- Tremendous performance increase of single threaded micro processors
- Two key effects:
 - Number of transistors on a chip doubled every 18 months (Moore's Law [12])
 - Power density of transistors was constant
→ Smaller transistors required less voltage and current (Dennard scaling [3])
- DBMSs got faster automatically with faster processors

The Failure of Dennard Scaling


- Transistors got smaller and smaller
 - With smaller transistors the leakage current increases
 - Higher leakage current results in higher power consumption of transistors
- Constant chip size and an increasing number of transistors increases overall power consumption and the produced heat [9]

The Power Wall and the Multi Core Era

- Practical limits on amount of power a processor can use:
 1. Cooling
 2. Energy Consumption
- Consequence: modern processors are limited by a fixed energy budget (power wall) [1]
- In 2004, Intel canceled Tejas and Jayhawk processors
→ Since then, they are building multi-core CPUs!

The Free Lunch is Over

Herb Sutter: Dr. Dobb's Journal 30(3), 2005



<http://www.gotw.ca/publications/concurrency-ddj.htm>

Limitations of Scaling The Number of Cores

- Since 2005, vendors concentrated on processors with multiple cores to use the increasing number of transistors to increase performance
- Benefit of increasing parallelism will not justify the costs of process scaling (i.e., creating even smaller transistors) [4]
- Experts predict that this trend will not scale beyond a few hundred cores [5]
- Examples:
 - Multi-core designs at 22 nm require 21 % of a chip's transistors to be powered off
 - With 8 nm transistors, it will be 50 %! [4]

Dark Silicon and the Heterogeneous Many Core Age

To remain in the power constraints, modern processors can either:

- Operate with lower clock rate or
- Turn off parts of the chip → *dark silicon* [1, 4, 5]

How can we exploit the increasing number of transistors?

- Provide a large number of cores that are *specialized* for certain tasks
- Cores that are most suitable for the current workload are powered on until the energy budget is reached [4, 5]
- Inevitably, this will cause processors to become increasingly *heterogeneous*

Today's and Tomorrow's Processors

- Future machines are expected to consist of a set of heterogeneous processors
- Each processor is optimized for a certain application scenario [1, 5]

This trend has already become commodity in the form of:

- Graphics processing units (GPUs)
- Many integrated cores architectures (MICs)
- Field-programmable gate arrays (FPGAs)
- Accelerated processing units (APUs)
- Processor vendors also combine heterogeneous processors on a single die (e.g., APUs)

A NEW ERA OF PROCESSOR PERFORMANCE

Single-Core Era

Enabled by:

- ✓ Moore's Law
- ✓ Voltage Scaling

Constrained by:

- ✗ Power
- ✗ Complexity

Assembly → C/C++ → Java ...



Multi-Core Era


Enabled by:

- ✓ Moore's Law
- ✓ SMP architecture

Constrained by:

- ✗ Power
- ✗ Parallel SW
- ✗ Scalability

pthreads → OpenMP / TBB ...



Heterogeneous Systems Era


Enabled by:

- ✓ Abundant data parallelism
- ✓ Power efficient GPUs

Temporarily Constrained by:

- ✗ Programming models
- ✗ Comm.overhead

Shader → CUDA → OpenCL → !!!



Databases in the Heterogeneous Many Core Age

- Not taking advantage of heterogeneous processors for query processing leaves available resources unused
- On the contrary, we should use these processors to accelerate query processing, e.g., efficient database algorithms:
 - GPUs [6]
 - APUs [7, 8]
 - MICs [10, 11]
 - FPGAs [2, 13]
- We need to think about *how* we can exploit all processors during query processing

Outline

Motivation

GPU Architecture

Relational Co-Processing on GPUs

Query Processing & Data Transfer Bottleneck

Summary

GPGPU for DBMS?

- tremendous computation power & massive parallelism

Processor	FLOPS	Cores
Intel Core i7 (Sandy Bridge)	≈ 102 GFLOPS	4
Tesla K40 (Kepler)	≈ 1.66 TFLOPS	2880

improve performance of


- main-memory DBMS
- compute-intensive tasks: expensive predicates (e.g. for spatial data), UDFs (statistical functions, data mining), query optimization, ...

General-Purpose GPUs (GPGPUs)

Original GPU design based on graphics pipeline not flexible enough.

- geometry shaders idle for pixel-heavy workloads and vice versa
- **unified model** with general-purpose cores

Thus: Design inspired by CPUs, but different



CPU




GPU

Rationale: Optimize for **throughput**, not for **latency**.

Example: NVIDIA GPUs


NVIDIA GTX 280



source: www.hardwaresecrets.com

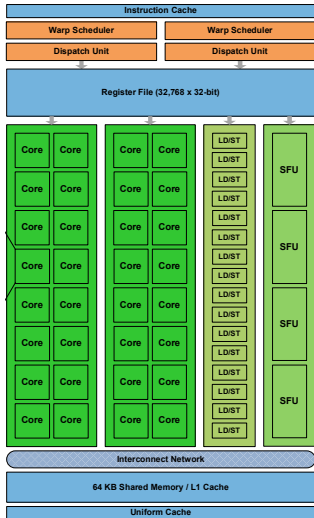
- 10 Thread Processing Clusters
- 10×3 Streaming Multiprocessors
- $10 \times 3 \times 8$ Scalar Processor Cores
 - More like ALUs (↗ slide 14)
- Each Multiprocessor:
 - 16k 32-bit registers
 - 16 kB shared memory
 - up to 1024 threads (may be limited by registers and/or memory)

Inside a Streaming Multiprocessor



- 8 Scalar Processors (Thread Processors)
 - single-precision floating point
 - 32-bit and 64-bit integer
- 2 Special Function Units
 - sin, cos, log, exp
- Double Precision unit
- 16 kB Shared Memory
- Each Streaming Multiprocessor: up to 1,024 threads.
- GTX 280: 30 Streaming Multiprocessors
 - 30,720 concurrent threads (!)

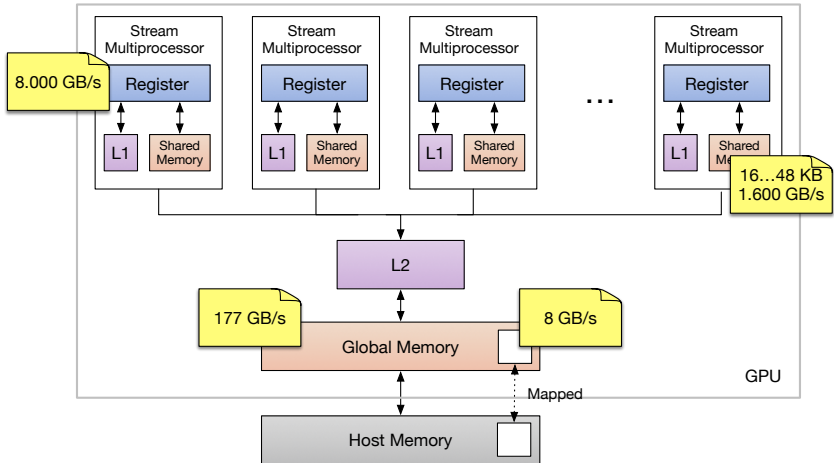
Inside a Streaming Multiprocessor: NVIDIA Fermi



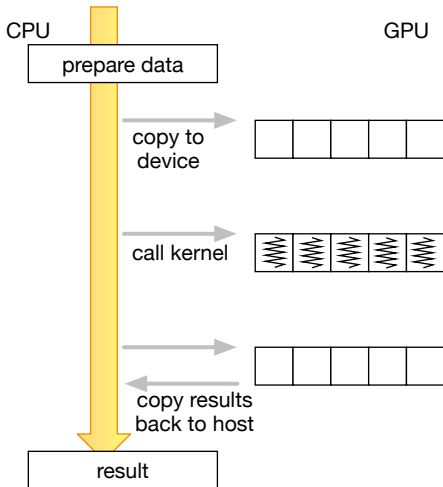
- 32 “cores” (thread processors) per streaming multiprocessor (SM)
- but fewer SMs per GPU: 16 (vs. 30 in GT200 architecture)
- 512 “cores” total
- “cores” now double-precision-capable

Source: NVIDIA Fermi White Paper

GPU: Memory Model



GPU: Execution Model



- data copying host (CPU) \leftrightarrow device (GPU)
- invocations of **compute kernels**
- kernels run asynchronously

Threads on a GPU

To handle 10,000s of threads efficiently, keep things simple.

- Don't try to **reduce** latency, but **hide** it.
 - **Large thread pool** rather than caches
(This idea is similar to SMT in commodity CPUs)
- Assume **data parallelism** and restrict **synchronization**.
 - Threads and small **groups** of threads use local memories.
 - Synchronization only within those groups (more later).
- Hardware **thread scheduling** (simple, in-order).
 - Schedule threads in **batches** (\leadsto “warps”).

Scheduling in Batches

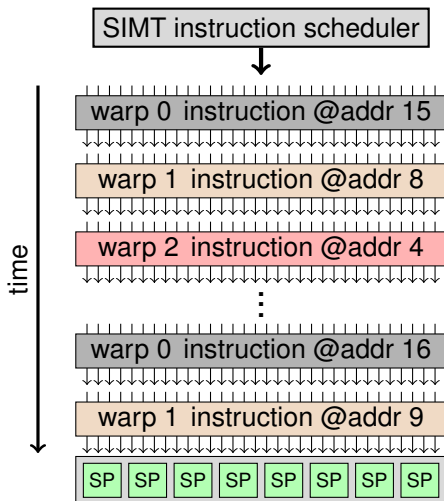
- In SM threads are scheduled in units of 32, called **warps**.
- **Warp**: Set of 32 parallel threads that start together at the same program address.



warp (dt. Kett- oder Längsfaden)

- Warps are scheduled with zero-overhead
- Scoreboard is used to track which warps are ready to execute
- GTX 280: 32 warps per multiprocessor (1024 threads)
- newer cards: 48 warps per multiprocessor (1536 threads)

SPMD / SIMT Processing



- **SIMT**: Single Instruction, Multiple Threads
- All threads execute the same instruction.
- Threads are split into warps by increasing thread IDs (warp 0 contains thread 0).
- At each time step scheduler selects warp ready to execute (*i.e.*, all its data are available)
- NVIDIA Fermi: dual issue; issue two warps at once

Warps and Latency Hiding

Some runtime characteristics:

- Issuing a warp instruction takes **4 cycles** (8 scalar processors).
- Register write-read latency: **24 cycles**.
- Global (off-chip) memory access: \approx **400 cycles**.

Threads are executed **in-order**.

- **Hide latencies** by executing other warps when one is paused.
- Need **enough warps** to fully hide latency.


e.g.,

- Need $24/4 = 6$ warps to hide register dependency latency.
- Need $400/4 = 100$ instructions to hide memory access latency.
If every 8th instruction is a memory access, $100/8 \approx 13$ warps would be enough.

CPUs vs. GPUs


CPU: task parallelism

- relatively heavyweight threads
- 10s of threads on 10s of cores
- each thread managed explicitly
- threads run different code



GPU: data parallelism

- lightweight threads
 - 10,000s of threads on 100s of cores
 - threads scheduled in batches
 - all threads run same code
- SPMD, single program, multiple data




APU: The next step?

CPU + GPU

- GPU kernels are fast
- but GPU has only PCIe bus as I/O path
- PCIe 2.1 / 16 lanes: theoretical throughput of 8 GB/s, but lower in practice \leadsto **data transfer can overwhelm runtime of kernels**

APU = Accelerated Processing Unit, e.g. AMD A10



- Example: A10-7850K = 12 cores (4 CPU + 6 GPU = 512 execution units), L2 Cache = 4 MB, 32 GB RAM
- CPU and GPU are integrated in the same chip
- ... and share the L2 cache

Outline

Motivation

GPU Architecture

Relational Co-Processing on GPUs

Query Processing & Data Transfer Bottleneck

Summary

Relational Co-Processing on GPUs

- Selection
- Join
- Aggregation

Relational Co-Processing on GPUs


- Selection
- Join
- Aggregation

We focus on selections here.

Selection

Choose a subset of tuples from a relation R satisfying a predicate and discard the rest:

pred: $val < 5$



```
algorithm:  
unsigned int i=0;  
for(i=0;i<n;i++){  
    if(pred(val[i]))  
        res.add(val[i]);  
}
```

Parallel Selections

How can we parallelize selections efficiently?

- Concurrent writes may corrupt data structures
→ Ensure correctness by latches
- Latching may serialize threads and nullify the performance gained by parallel execution
→ Need to ensure correctness without latching
- Key Idea: Pre-compute write locations


B. He et. al. Relational query coprocessing on graphics processors. TODS, 34(4):21, 2009.

Prefix Scans

- Important building block for parallel programs
- Applies a binary operator to an array
- Example: prefix sum
- Given an input array R_{in} , R_{out} is computed as follows:
$$R_{out}[i] = R_{in}[0] + \dots + R_{in}[i-1] \quad (1 \leq i < |R_{in}|)$$
$$R_{out}[0] = 0$$

B. He et. al. Relational query coprocessing on graphics processors. TODS, 34(4):21, 2009.

Example: Prefix Sum



Parallel Filter

- Create an array *flags* of the same size as *R* and init with zeros
- For each tuple set corresponding flag in *flags* if and only if the current tuple matches the predicate
 - *flags* array contains a 1 if the corresponding tuple in *R* is part of the result
 - The sum of the values in *flags* is the number of result tuples $\#rt$
- Compute the prefix sum of *flags* and store it in array *ps*
 - Now we have the write locations for each tuple in the result buffer

B. He et. al. Relational query coprocessing on graphics processors. TODS, 34(4):21, 2009.

Parallel Filter

- Create the result buffer *res* of size $\#rt$
- Scan *flags*: if(flags[i]==1) write R[i] to position ps[i] in the result buffer:


do in parallel:

```
for (unsigned int i=0; i<n; i++){  
    if ( flags [ i ]==1){  
        unsigned int res_write_index=ps [ i ];  
        res [ res_write_index ]=R [ i ];  
    }  
}
```

B. He et. al. Relational query coprocessing on graphics processors. TODS, 34(4):21, 2009.

Parallel Filter: Example

predicate: $val < 5$



build flag array:
 $\text{if}(\text{pred}(\text{val}[i]))$
 $\text{flags}[i]=1;$
else
 $\text{flags}[i]=0;$

compute
prefix sum ps
from $flags$

scan $flags$ and
write $\text{val}[i]$ to
position $ps[i]$
in result array

B. He et. al. Relational query coprocessing on graphics processors. TODS, 34(4):21, 2009.

Outline

Motivation


GPU Architecture

Relational Co-Processing on GPUs

Query Processing & Data Transfer Bottleneck

Summary

Communication Bottleneck




Resource Limitations

- Co-Processor's memory capacity quite small (up to 16 GB)
→ Cannot fit all data on the co-processor
- Ad-hoc data transfers are expensive
→ Cache input data in co-processor memory and process data locally as much as possible

S. Breß, H. Funke, J. Teubner. Robust Query Processing in Co-Processor-accelerated Databases. In SIGMOD, 2016.

Cache input data for best performance



S. Breß, H. Funke, J. Teubner. Robust Query Processing in Co-Processor-accelerated Databases. In SIGMOD, 2016.

Cache Thrashing

Background:

- Assume a fixed set of re-occurring queries (for discussion only)
- When a query accesses data, it is cached on the co-processor
- If cache is full, data from other queries is evicted

Effects:

- Occurs when the working set does not fit into data cache of a co-processor
- Queries evict each others data from the cache
- In the worst case, we never have a cache hit (if every query accesses different data)
 - We pay the full data transfer cost for every query despite having a cache!

S. Breß. H. Funke, J. Teubner. Robust Query Processing in Co-Processor-accelerated Databases. In SIGMOD, 2016.

Operator Placement and Data Placement

Operator Placement:

- Select a processor for each operator in the query plan
- Aims to increase data locality in a single query

Data Placement:

- Select which part of the database is stored in a co-processor's memory
→ Reserve part of the co-processors device memory as data cache
- Aims to increase data locality across queries
- Similar to classical page replacement strategies in disk-based databases

S. Breß, H. Funke, J. Teubner. Robust Query Processing in Co-Processor-accelerated Databases. In SIGMOD, 2016.


Duality of Operator and Data Placement

We can either

- Decide on an operator placement, and transfer the input data accordingly, or
- Decide on a data placement, and place the operators accordingly

S. Breß, H. Funke, J. Teubner. Robust Query Processing in Co-Processor-accelerated Databases. In SIGMOD, 2016.

Classical Operator-Driven Data Placement



Data-Placement-Driven Operator Placement

Key Observations:

1. We do not have to use a co-processor, if we expect high overheads for data transfer
→ Process data on CPU, and use co-processor only when advantageous (e.g., input data cached)
2. If the input data for an operator is cached on the co-processor, the co-processor is (most of the time) faster than the CPU
→ Push operators (code) to data

S. Breß, H. Funke, J. Teubner. Robust Query Processing in Co-Processor-accelerated Databases. In SIGMOD, 2016.


Data-Placement-Driven Operator Placement

Algorithm consists of two stages:


1. Periodically run a data placement thread in the background:
 - Identify part of the working set that fits in co-processors data cache
→ Use most recently or most frequently used columns/data pages
 - Transfer the selected columns/data pages to the co-processors cache
2. For each operator:
 - Test whether input data is cached on the co-processor (this includes result data from previous operators)
 - If the data is cached, execute operator on co-processor, if not, execute operator on CPU

S. Breß, H. Funke, J. Teubner. Robust Query Processing in Co-Processor-accelerated Databases. In SIGMOD, 2016.

Data-Placement-Driven Operator Placement




Data to Query or Query to Data?



Data-driven operator placement solves the cache thrashing problem

Data to Query or Query to Data? (2)



Summary

Cache Thrashing

- Problem: Scarce cache memory leads to expensive eviction and recaching and slows down query processing
- Solution: Data-centric operator placement, which pushes operators to processors that have data cached

S. Breß, H. Funke, J. Teubner. Robust Query Processing in Co-Processor-accelerated Databases. In SIGMOD, 2016.



<http://cogadb.dfki.de>

Outline

Motivation

GPU Architecture

Relational Co-Processing on GPUs

Query Processing & Data Transfer Bottleneck

Summary

Summary

Hardware trends

- Increasing heterogeneity: CPUs, GPUs, co-processors
- Many-core processors, even on CPUs
- Disk IO is not the limiting factor anymore

Consequences

- Need to actively adjust the DBMS to make use of heterogeneous processors (e.g., relational operators, query processor)
- Requires to rethink architectural aspects of the DBMS (e.g., columns stores, compression)

Processors are parallel and heterogenous, we have to deal with it!

References I

- [1] S. Borkar and A. A. Chien.
The future of microprocessors.
Communications of the ACM, 54(5):67–77, 2011.
- [2] J. Casper and K. Olukotun.
Hardware acceleration of database operations.
In *FPGA*, pages 151–160. ACM, 2014.
- [3] R. Dennard, V. Rideout, E. Bassous, and A. LeBlanc.
Design of ion-implanted mosfet's with very small physical dimensions.
IEEE Journal of Solid-State Circuits, 9(5):256–268, 1974.
- [4] H. Esmailzadeh, E. Blem, R. St. Amant, K. Sankaralingam, and D. Burger.
Dark silicon and the end of multicore scaling.
In *ISCA*, pages 365–376. ACM, 2011.
- [5] N. Hardavellas, M. Ferdman, B. Falsafi, and A. Ailamaki.
Toward dark silicon in servers.
IEEE Micro, 31(4):6–15, 2011.
- [6] B. He, M. Lu, K. Yang, R. Fang, N. K. Govindaraju, Q. Luo, and P. V. Sander.
Relational query co-processing on graphics processors.
In *ACM Trans. Database Syst.*, volume 34. ACM, 2009.

References II

- [7] J. He, M. Lu, and B. He.
Revisiting co-processing for hash joins on the coupled CPU-GPU architecture.
Proc. VLDB Endow., 6(10):889–900, 2013.
- [8] J. He, S. Zhang, and B. He.
In-cache query co-processing on coupled CPU-GPU architectures.
Proc. VLDB Endow., 8(4):329–340, 2014.
- [9] J. L. Hennessy and D. A. Patterson.
Computer Architecture: A Quantitative Approach.
Morgan Kaufmann Publishers Inc., 5th edition, 2011.
- [10] S. Jha, B. He, M. Lu, X. Cheng, and H. P. Huynh.
Improving main memory hash joins on intel xeon phi processors: An experimental approach.
PVLDB, 8(6):642–653, 2015.
- [11] M. Lu, L. Zhang, H. P. Huynh, Z. Ong, Y. Liang, B. He, R. Goh, and R. Huynh.
Optimizing the mapreduce framework on Intel Xeon Phi coprocessor.
In *Big Data*, pages 125–130. IEEE, 2013.
- [12] G. Moore.
Cramming more components onto integrated circuits.
Electronics, 38(8), 1965.

References III

- [13] R. Mueller, J. Teubner, and G. Alonso.
Data processing on FPGAs.
PVLDB, 2(1):910–921, 2009.
- [14] K. Sattler, J. Teubner, F. Beier, and S. Breß.
Many-core-architekturen zur datenbankbeschleunigung.
In *Datenbanksysteme für Business, Technologie und Web (BTW 2015) - Workshopband*,
2.-3. März 2015, Hamburg, Germany, pages 269–270, 2015.